

AI ASSISTED CODING - 17.3

Baisa Vaishnavi

2303A52142

Batch - 41

Task 1 Employee Dataset Cleaning

Task:

Use AI assistance to clean an employee information dataset.


Instructions:

- Handle missing values in salary, designation, and experience
- Remove duplicate employee IDs
- Standardize text values (job titles in lowercase)
- Encode categorical features (designation, employment_type)

Prompt:

Employee Dataset Cleaning . Use AI assistance to clean an employee information dataset. Instructions: • Handle missing values in salary, designation, and experience • Remove duplicate employee IDs • Standardize text values (job titles in lowercase) • Encode categorical features (designation, employment_type) Expected Output: A clean dataset suitable for HR analytics and prediction tasks. Give me code for this dataset and give me proper information of the employees and and give me some invalid employee information.

Code & output:



Subtask:

Extract the contents of the `/content/archive.zip` file to access the employee dataset.

Reasoning: To extract the contents of the `/content/archive.zip` file as per the subtask, I will use the `zipfile` module in a Python code block.

```
[1]
✓ 0s
import zipfile
import os

zip_file_path = '/content/archive.zip'
extract_dir = '/content/'


# Create the extraction directory if it doesn't exist
os.makedirs(extract_dir, exist_ok=True)

# Open the zip file and extract its contents
with zipfile.ZipFile(zip_file_path, 'r') as zip_ref:
    zip_ref.extractall(extract_dir)

print(f'Archive '{zip_file_path}' extracted to '{extract_dir}' successfully.
```

... Archive '/content/archive.zip' extracted to '/content/' successfully.

Load Employee Dataset



Reasoning: I will load the 'Employee.csv' file into a pandas DataFrame and display its first 5 rows as per the instructions to identify the dataset structure.

```
[2]
✓ 1s
import pandas as pd

file_path = '/content/Employee.csv'
df_employee = pd.read_csv(file_path)

print(f"Successfully loaded '{file_path}' into df_employee.")
print("First 5 rows of df_employee:")
print(df_employee.head())
```

Successfully loaded '/content/Employee.csv' into df_employee.
First 5 rows of df_employee:

	Education	JoiningYear	City	PaymentTier	Age	Gender	EverBenched
0	Bachelors	2017	Bangalore	3	34	Male	No
1	Bachelors	2013	Pune	1	28	Female	No
2	Bachelors	2014	New Delhi	3	38	Female	No
3	Masters	2016	Bangalore	3	27	Male	No
4	Masters	2017	Pune	3	24	Male	Yes

	ExperienceInCurrentDomain	LeaveOrNot
0	0	0
1	3	1
2	2	0
3	5	1
4	2	1

Handle Missing Values

[3]
✓ Os

```
print("Columns in df_employee:")
print(df_employee.columns)

print("\nCalculating missing values for all columns:")
missing_values_before = df_employee.isnull().sum()
print(missing_values_before)

# Assuming 'ExperienceInCurrentDomain' is the 'experience' column mentioned in the subtask
# Check if 'ExperienceInCurrentDomain' has missing values and impute if necessary
if 'ExperienceInCurrentDomain' in df_employee.columns and missing_values_before['ExperienceInCurrentDomain'] > 0:
    median_experience = df_employee['ExperienceInCurrentDomain'].median()
    df_employee['ExperienceInCurrentDomain'].fillna(median_experience, inplace=True)
    print(f"\nImputed missing values in 'ExperienceInCurrentDomain' with its median: {median_experience}")

print("\nMissing values after imputation:")
print(df_employee.isnull().sum())

# Explicitly state that 'salary' and 'designation' columns were not found
if 'salary' not in df_employee.columns:
    print("\nNote: The 'salary' column was not found in the dataset, so no missing values were handled for it.")
if 'designation' not in df_employee.columns:
    print("\nNote: The 'designation' column was not found in the dataset, so no missing values were handled for it.")
```

```
Columns in df_employee:
Index(['Education', 'JoiningYear', 'City', 'PaymentTier', 'Age', 'Gender',
      'EverBenched', 'ExperienceInCurrentDomain', 'LeaveOrNot'],
      dtype='object')
```

Calculating missing values for all columns:

Education	0
JoiningYear	0
City	0
PaymentTier	0
Age	0
Gender	0
EverBenched	0
ExperienceInCurrentDomain	0
LeaveOrNot	0

dtype: int64

Missing values after imputation:

Education	0
JoiningYear	0
City	0
PaymentTier	0
Age	0
Gender	0
EverBenched	0
ExperienceInCurrentDomain	0
LeaveOrNot	0

dtype: int64

Note: The 'salary' column was not found in the dataset, so no missing values were handled for it.

Note: The 'designation' column was not found in the dataset, so no missing values were handled for it.

Remove Duplicate Employee IDs

```
[4]
✓ Os
print("Original data types:")
print(df_employee.dtypes)

# Identify columns with 'object' data type (typically text fields)
object_columns = df_employee.select_dtypes(include='object').columns

print(f"\nIdentified object columns for standardization: {list(object_columns)}")

# Convert string values in object columns to lowercase
for col in object_columns:
    # Ensure values are strings before applying .str.lower()
    df_employee[col] = df_employee[col].astype(str).str.lower()

print("\nText standardization complete. Displaying first 5 rows with standardized text:")
print(df_employee.head())

# Note about 'designation' and 'job titles' as per instructions
if 'designation' not in df_employee.columns and 'job titles' not in df_employee.columns:
    print("\nNote: The columns 'designation' and 'job titles' mentioned in the subtask were not found in the dataset. Other relevant
```

```
✓ Os
✓ ...
Original data types:
Education      object
JoiningYear    int64
City           object
PaymentTier    int64
Age            int64
Gender         object
EverBenched    object
ExperienceInCurrentDomain  int64
LeaveOrNot      int64
dtype: object

Identified object columns for standardization: ['Education', 'City', 'Gender', 'EverBenched']

Text standardization complete. Displaying first 5 rows with standardized text:
  Education  JoiningYear  City  PaymentTier  Age  Gender  EverBenched  \
0  bachelors    2017    bangalore        3   34    male         no
1  bachelors    2013     pune        1   28    female        no
2  bachelors    2014  new delhi        3   38    female        no
3  masters     2016    bangalore        3   27    male         no
4  masters     2017     pune        3   24    male         yes

  ExperienceInCurrentDomain  LeaveOrNot
0                        0            0
1                        3            1
2                        2            0
3                        5            1
4                        2            1

Note: The columns 'designation' and 'job titles' mentioned in the subtask were not found in the dataset. Other relevant text
```

Encode Categorical Features

```

import pandas as pd
# Identify categorical columns for encoding
categorical_cols = ['Education', 'City', 'Gender', 'EverBenched']

print(f"Categorical columns to be one-hot encoded: {categorical_cols}")

# Apply one-hot encoding
df_employee = pd.get_dummies(df_employee, columns=categorical_cols, drop_first=False)

print("\nOne-hot encoding complete. Displaying first 5 rows of the updated DataFrame:")
print(df_employee.head())

print("\nData types of the updated DataFrame:")
print(df_employee.dtypes)

```

... Categorical columns to be one-hot encoded: ['Education', 'City', 'Gender', 'EverBenched']

One-hot encoding complete. Displaying first 5 rows of the updated DataFrame:

	JoiningYear	PaymentTier	Age	ExperienceInCurrentDomain	LeaveOrNot	\
0	2017	3	34		0	0
1	2013	1	28		3	1
2	2014	3	38		2	0
3	2016	3	27		5	1
4	2017	3	24		2	1

	Education_bachelors	Education_masters	Education_phd	City_bangalore	\
0	True	False	False	True	
1	True	False	False	False	
2	True	False	False	False	
3	False	True	False	True	
4	False	True	False	False	

... City_new delhi City_pune Gender_female Gender_male EverBenched_no \

0	False	False	False	True	True
1	False	True	True	False	True
2	True	False	True	False	True
3	False	False	False	True	True
4	False	True	False	True	False

	EverBenched_yes
0	False
1	False
2	False
3	False
4	True

Data types of the updated DataFrame:

JoiningYear	int64
PaymentTier	int64
Age	int64
ExperienceInCurrentDomain	int64
LeaveOrNot	int64
Education_bachelors	bool
Education_masters	bool
Education_phd	bool
City_bangalore	bool
City_new delhi	bool
City_pune	bool
Gender_female	bool
Gender_male	bool
EverBenched_no	bool
EverBenched_yes	bool
dtype:	object

Identify Invalid Employee Information

[6]
✓ Os

```
print("Descriptive statistics for numerical columns:")
print(df_employee.describe())

# Identify potential invalid or outlier information based on common sense

# 1. Check for unrealistic 'Age' values (e.g., less than 18)
invalid_age = df_employee[df_employee['Age'] < 18]
if not invalid_age.empty:
    print("\n--- Potentially Invalid Age (less than 18) ---")
    print(invalid_age)
else:
    print("\nNo employees found with age less than 18.")

# 2. Check for unrealistic 'JoiningYear' (e.g., in the future or too far in the past)
# Assuming the current year for comparison. Current year is 2024
current_year = 2024
invalid_joining_year_future = df_employee[df_employee['JoiningYear'] > current_year]
invalid_joining_year_past = df_employee[df_employee['JoiningYear'] < 1950] # Arbitrary past limit

if not invalid_joining_year_future.empty:
    print(f"\n--- Potentially Invalid JoiningYear (after {current_year}) ---")
    print(invalid_joining_year_future)
else:
    print(f"\nNo employees found with joining year after {current_year}.")

if not invalid_joining_year_past.empty:
    print(f"\n--- Potentially Invalid JoiningYear (before 1950) ---")
    print(invalid_joining_year_past)
```

[6]
✓ Os

```
else:
    print("\nNo employees found with joining year before 1950.")

# 3. Check for 'ExperienceInCurrentDomain' values (e.g., negative or greater than age)
invalid_experience_negative = df_employee[df_employee['ExperienceInCurrentDomain'] < 0]
if not invalid_experience_negative.empty:
    print("\n--- Potentially Invalid ExperienceInCurrentDomain (negative) ---")
    print(invalid_experience_negative)
else:
    print("\nNo employees found with negative experience in current domain.")

invalid_experience_greater_than_age = df_employee[df_employee['ExperienceInCurrentDomain'] >= df_employee['Age']]
if not invalid_experience_greater_than_age.empty:
    print("\n--- Potentially Invalid ExperienceInCurrentDomain (greater than or equal to Age) ---")
    print(invalid_experience_greater_than_age)
else:
    print("\nNo employees found with experience in current domain greater than or equal to their age.")

print("\nFinished checking for invalid employee information.")
```

Descriptive statistics for numerical columns:

	JoiningYear	PaymentTier	Age	ExperienceInCurrentDomain \
count	4653.000000	4653.000000	4653.000000	4653.000000
mean	2015.062970	2.698259	29.393295	2.905652
std	1.863377	0.561435	4.826087	1.558240
min	2012.000000	1.000000	22.000000	0.000000
25%	2013.000000	3.000000	26.000000	2.000000
50%	2015.000000	3.000000	28.000000	3.000000
75%	2017.000000	3.000000	32.000000	4.000000
max	2018.000000	3.000000	41.000000	7.000000

	LeaveOrNot
count	4653.000000
mean	0.343864
std	0.475047
min	0.000000
25%	0.000000
50%	0.000000
75%	1.000000
max	1.000000

No employees found with age less than 18.

No employees found with joining year after 2024.

No employees found with joining year before 1950.

No employees found with negative experience in current domain.

No employees found with experience in current domain greater than or equal to their age.

Finished checking for invalid employee information.

Display Cleaned Dataset Information

```
print("Shape of the cleaned df_employee DataFrame:")
print(df_employee.shape)

print("\nData types of df_employee DataFrame:")
print(df_employee.dtypes)

print("\nFirst 5 rows of the cleaned df_employee DataFrame:")
print(df_employee.head())
```

Shape of the cleaned df_employee DataFrame:
(4653, 15)

Data types of df_employee DataFrame:

```
JoiningYear      int64
PaymentTier      int64
Age              int64
ExperienceInCurrentDomain  int64
LeaveOrNot        int64
Education_bachelors    bool
Education_masters      bool
Education_phd          bool
City_bangalore         bool
City_new delhi         bool
City_pune              bool
Gender_female          bool
Gender_male            bool
EverBenched_no         bool
EverBenched_yes        bool
dtype: object
```

First 5 rows of the cleaned df_employee DataFrame:

	JoiningYear	PaymentTier	Age	ExperienceInCurrentDomain	LeaveOrNot	\
0	2017		3	34	0	0
1	2013		1	28	3	1
2	2014		3	38	2	0
3	2016		3	27	5	1
4	2017		3	24	2	1

	Education_bachelors	Education_masters	Education_phd	City_bangalore	\
0	True	False	False	True	
1	True	False	False	False	
2	True	False	False	False	
3	False	True	False	True	
4	False	True	False	False	

	City_new delhi	City_pune	Gender_female	Gender_male	EverBenched_no	\
0	False	False	False	True	True	
1	False	True	True	False	True	
2	True	False	True	False	True	
3	False	False	False	True	True	
4	False	True	False	True	False	

	EverBenched_yes
0	False
1	False
2	False
3	False
4	True

Justification:

- MissingValuesHandling
 - *Salary*→*Median*:Salarydistributionsareoftenskewed;
median is robust to outliers.

- *Designation* → *Mode*: Most frequent job title maintains categorical consistency.
- *Experience* → *Mean*: Numerical field; mean preserves overall distribution.
- Duplicate Removal
 - Removing duplicate `employee_id` ensures data integrity and avoids biased analytics.
- Text Standardization
 - Converting designation to lowercase eliminates case-sensitive inconsistencies (e.g., “Manager” vs “manager”).
- Categorical Encoding
 - Machine learning models require numerical inputs; encoding converts text into usable numeric form.

Task 2 – Student Performance Data Preprocessing

Task:

Preprocess student performance data using AI assistance.

Prompt:

Preprocess a student performance dataset.

Instructions:

- Convert `exam_date` column into datetime format
- Extract new features (month, semester, day)
- Normalize `total_marks`
- Detect and handle outliers in marks using IQR or Z-score

Code & Output

```
[3]  import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler

np.random.seed(0)
dates = pd.date_range(start="2024-01-01", periods=60)

df = pd.DataFrame({
    "student_id": range(1, 61),
    "exam_date": np.random.choice(dates, 60),
    "total_marks": np.random.randint(200, 1000, 60)
})

# Add outliers
df.loc[5, 'total_marks'] = 1500
df.loc[10, 'total_marks'] = 50

# Convert datetime
df['exam_date'] = pd.to_datetime(df['exam_date'])

# Feature Extraction
df['month'] = df['exam_date'].dt.month
df['day'] = df['exam_date'].dt.day
df['semester'] = df['month'].apply(lambda x: 1 if x <= 6 else 2)

# Normalize
scaler = MinMaxScaler()
df['total_marks'] = scaler.fit_transform(df[['total_marks']])

# Remove Outliers (IQR)
Q1 = df['total_marks'].quantile(0.25)
Q3 = df['total_marks'].quantile(0.75)
IQR = Q3 - Q1

df = df[(df['total_marks'] >= Q1 - 1.5*IQR) &
        (df['total_marks'] <= Q3 + 1.5*IQR)]

print(df.head())
```

```
...  student_id  exam_date  total_marks  month  day  semester
0      1  2024-02-14      0.229655      2   14          1
1      2  2024-02-17      0.122759      2   17          1
2      3  2024-02-23      0.191724      2   23          1
3      4  2024-01-01      0.191724      1    1          1
4      5  2024-01-04      0.140000      1    4          1
```

Justification:

Datetime conversion → enables time-based analysis

Feature extraction → captures patterns (semester trends)

MinMax scaling → brings values to 0–1 range

Outlier removal → improves model accuracy

Task 3 – Student Health Data Preparation

Task:

Clean and preprocess student health records using AI scripts.

Prompt:

Prepare a student health dataset for machine learning.

Instructions:

- Handle missing values in height, weight, BMI
- Encode categorical features (medical_condition, fitness_status)
- Scale numerical features using standard scaling
- Split dataset into training and testing sets

Code&output:

```
[4] ✓ 0s
import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split

np.random.seed(1)

df = pd.DataFrame({
    "student_id": range(1, 51),
    "height": np.random.choice([150, 160, 170, 180, np.nan], 50),
    "weight": np.random.choice([50, 60, 70, 80, np.nan], 50),
    "BMI": np.random.choice([18, 22, 25, 30, np.nan], 50),
    "medical_condition": np.random.choice(["None", "Asthma", "Diabetes"], 50),
    "fitness_status": np.random.choice(["Fit", "Unfit"], 50)
})

# Handle missing values
df.fillna(df.mean(numeric_only=True), inplace=True)

# Encoding
le1, le2 = LabelEncoder(), LabelEncoder()
df['medical_condition'] = le1.fit_transform(df['medical_condition'])
df['fitness_status'] = le2.fit_transform(df['fitness_status'])

# Scaling
scaler = StandardScaler()
df[['height', 'weight', 'BMI']] = scaler.fit_transform(df[['height', 'weight', 'BMI']])
```

```
[4] ✓ 0s # Split
X = df.drop('fitness_status', axis=1)
y = df['fitness_status']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

print(X_train.shape, X_test.shape)
```

... (40, 5) (10, 5)

Justification:

Mean imputation → preserves dataset size

Encoding → converts categorical to numeric

Standard scaling → required for ML algorithms

Train-test split → evaluates model performance

Task 4: Real-Time Application: Ride-Sharing Data Cleaning

Scenario:

A ride-sharing company has trip data with errors.

Prompt:

Clean a ride-sharing dataset.

Requirements:

- Standardize pickup and drop locations
- Fill missing fare amounts with median values
- Remove duplicate ride entries
- Normalize driver ratings
- Generate a summary of data improvements

Code:

[2]
✓ 1s

```
import pandas as pd
import numpy as np
from sklearn.preprocessing import MinMaxScaler

#----SampleDataset----
np.random.seed(2)

locations = ["Chennai", "Adyar", "T Nagar", "Velachery", "Tambaram"]

df = pd.DataFrame({
    "ride_id": np.random.randint(1, 40, 60),
    "pickup": np.random.choice(locations, 60),
    "drop": np.random.choice(locations, 60),
    "fare": np.random.choice([100, 150, 200, 250, np.nan], 60),
    "driver_rating": np.random.uniform(3.0, 5.0, 60)
})

# Introduce inconsistencies
df.loc[0, 'pickup'] = "chennai"
df.loc[1, 'drop'] = "adyar"

# ---- 1. Standardize Locations ----
df['pickup'] = df['pickup'].str.lower()
df['drop'] = df['drop'].str.lower()

# ---- 2. Handle Missing Fare ----
df['fare'] = df['fare'].fillna(df['fare'].median())
```



```
#----3.RemoveDuplicateRides----
before_count = len(df)
df=df.drop_duplicates(subset='ride_id')
after_count = len(df)

#----4.NormalizeDriverRatings----
scaler = MinMaxScaler()
df['driver_rating'] = scaler.fit_transform(df[['driver_rating']])

# ---- 5. Summary ----
print("CLEANEDDATASETSAMPLE:")
print(df.head())

print("\nSUMMARY:")
print("Records Before Cleaning:", before_count)
print("RecordsAfterRemovingDuplicates:", after_count)
print("Missing Values:\n", df.isnull().sum())
```

Output:

```
... CLEANEDDATASETSAMPLE:
   ride_id  pickup      drop  fare  driver_rating
0       16  chennai  chennai  200.0         0.976876
1        9   t nagar   adyar  100.0         0.787486
2       23  tambaram  chennai  100.0         0.423726
3       19   t nagar  velachery 250.0         0.341723
4       12  velachery   adyar  150.0         0.539948

SUMMARY:
Records Before Cleaning: 60
RecordsAfterRemovingDuplicates: 28
Missing Values:
  ride_id      0
 pickup      0
  drop       0
  fare       0
 driver_rating 0
dtype: int64
```

Justification:

Location standardization → avoids duplication issues

Median for fare → robust to extreme values

Duplicate removal → ensures accurate records

Normalization → fair comparison of ratings

Summary → verifies cleaning effectiveness